

CS 33 – Discussion 1A

Computer System Organization

Week 3

Questions before we start?

Logistics

- Lab 1 - Datalab due tonight 11:59pm
- Lab 2 – Bomb lab released tomorrow – Due Friday May 1st 11:59
 - Probably have HW 2 assigned and due before then
- Please submit suggestions on CCLE TA-site on how to improve OH or DISC

Agenda

10am-11am PST:

- x86-64 basics continued
 - Condition codes
 - Conditional Branching
 - Switch statements/Jump tables
 - Structs/Data alignment
 - Register saving conventions
- Recursion/Tail Recursion code example

11am-11:50am PST:

- LA worksheet

Agenda

If time permits:

- Detail Bomblab
- Answer questions on Datalab

x86-64 Basics – Condition codes

- Condition codes are set implicitly by arithmetic instructions
- Explicitly set by setX instructions or cmp/test instruction





Instruction		Description	Page #
cmp	S_2, S_1	Set condition codes according to $S_1 - S_2$	185
test	S_2, S_1	Set condition codes according to S_1 & S_2	185

- Can be thought of as single bit registers

x86-64 Basics – Condition codes

Condition Codes (Implicit Setting)

Single bit registers

 CF	Carry Flag (for unsigned)	 SF	Sign Flag (for signed)
 ZF	Zero Flag	 OF	Overflow Flag (for signed)

Implicitly set (think of it as side effect) by arithmetic operations

Example: `addq Src, Dest` \leftrightarrow `t = a+b`

CF set if carry out from most significant bit (unsigned overflow)

ZF set if `t == 0`

SF set if `t < 0` (as signed)

OF set if two's-complement (signed) overflow





`(a>0 && b>0 && t<0) || (a<0 && b<0 && t>=0)`

Not set by `leaq` instruction

x86-64 Basics – Condition codes

Condition Codes (Implicit Setting)

Single bit registers

 CF	Carry Flag (for unsigned)	 SF	Sign Flag (for signed)
 ZF	Zero Flag	 OF	Overflow Flag (for signed)

Implicitly set (think of it as side effect) by arithmetic operations

Example: `addq Src, Dest` \leftrightarrow `t = a+b`

CF set if carry out from most significant bit (unsigned overflow)

ZF set if `t == 0`

SF set if `t < 0` (as signed)

OF set if two's-complement (signed) overflow

`(a>0 && b>0 && t<0) || (a<0 && b<0 && t>=0)`

Not set by `leaq` instruction

x86-64 Basics – Condition codes

Condition Codes (Explicit Setting: Compare)

🌀 Explicit Setting by Compare Instruction

🌀 `cmpq Src2, Src1`

🌀 `cmpq b, a` like computing `a-b` without setting destination

🌀 **CF set** if carry out from most significant bit (used for unsigned comparisons)

🌀 **ZF set** if `a == b`

🌀 **SF set** if `(a-b) < 0` (as signed)

🌀 **OF set** if two's-complement (signed) overflow

`(a>0 && b<0 && (a-b)<0) || (a<0 && b>0 && (a-b)>0)`

x86-64 Basics – Condition codes

Condition Codes (Explicit Setting: Test)

Explicit Setting by Test instruction

testq *Src2*, *Src1*

testq *b*, *a* like computing *a & b* without setting destination

Sets condition codes based on value of *Src1* & *Src2*

Useful to have one of the operands be a mask

ZF set when *a & b* == 0

SF set when *a & b* < 0

x86-64 Basics – Condition codes

🔄 SetX Instructions

- 🔄 Set low-order byte of destination to 0 or 1 based on combinations of condition codes
- 🔄 Does not alter remaining 7 bytes

SetX	Condition	Description
<code>sete</code>	<code>ZF</code>	Equal / Zero
<code>setne</code>	<code>~ZF</code>	Not Equal / Not Zero
<code>sets</code>	<code>SF</code>	Negative
<code>setns</code>	<code>~SF</code>	Nonnegative
<code>setg</code>	<code>~ (SF^OF) & ~ZF</code>	Greater (Signed)
<code>setge</code>	<code>~ (SF^OF)</code>	Greater or Equal (Signed)
<code>setl</code>	<code>(SF^OF)</code>	Less (Signed)
<code>setle</code>	<code>(SF^OF) ZF</code>	Less or Equal (Signed)
<code>seta</code>	<code>~CF & ~ZF</code>	Above (unsigned)
<code>setb</code>	<code>CF</code>	Below (unsigned)

x86-64 Basics – Conditional Branching

3.4.2 Jump Instructions

Instruction		Description	Condition Code	Page #
jmp	<i>Label</i>	Jump to label		189
jmp	<i>*Operand</i>	Jump to specified location		189
jz / jbe	<i>Label</i>	Jump if equal/zero	ZF	189
jne / jnb	<i>Label</i>	Jump if not equal/nonzero	\sim ZF	189
js	<i>Label</i>	Jump if negative	SF	189
jns	<i>Label</i>	Jump if nonnegative	\sim SF	189
jg / jnle	<i>Label</i>	Jump if greater (signed)	\sim (SF \wedge OF) & \sim ZF	189
jge / jnl	<i>Label</i>	Jump if greater or equal (signed)	\sim (SF \wedge OF)	189
jl / jnge	<i>Label</i>	Jump if less (signed)	SF \wedge OF	189
jle / jng	<i>Label</i>	Jump if less or equal	(SF \wedge OF) ZF	189
ja / jnbe	<i>Label</i>	Jump if above (unsigned)	\sim CF & \sim ZF	189
jae / jnb	<i>Label</i>	Jump if above or equal (unsigned)	\sim CF	189
jb / jnae	<i>Label</i>	Jump if below (unsigned)	CF	189
jbe / jna	<i>Label</i>	Jump if below or equal (unsigned)	CF ZF	189

x86-64 Basics – Switch Statement

```
long switch_eg
(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
            /* Fall Through */
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}
```

Switch Statement Example

Multiple case labels

Here: 5 & 6

Fall through cases

Here: 2

Missing cases

Here: 4



x86-64 Basics – Switch Statement

Jump Table Structure

Switch Form

```
switch(x) {  
  case val_0:  
    Block 0  
  case val_1:  
    Block 1  
    . . .  
  case val_n-1:  
    Block n-1  
}
```

Jump Table

jtab:	Targ0
	Targ1
	Targ2
	⋮
	Targn-1

Jump Targets

Targ0: Code Block 0

Targ1: Code Block 1

Targ2: Code Block 2

⋮

Targn-1: Code Block n-1

Translation (Extended C)

```
goto *JTab[x];
```



x86-64 Basics – Switch Statement

Switch Statement Example

```
long switch_eg(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        . . .
    }
    return w;
}
```

Setup:

```
switch_eg:
    movq    %rdx, %rcx
    cmpq    $6, %rdi        # x:6
    ja      .L8              # Use default
    jmp     *.L4(, %rdi, 8)   # goto *JTab[x]
```

*Indirect
jump*



Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rdx	Argument z
%rax	Return value

Jump table


```
.section      .rodata
    .align 8
.L4:
    .quad     .L8    # x = 0
    .quad     .L3    # x = 1
    .quad     .L5    # x = 2
    .quad     .L9    # x = 3
    .quad     .L8    # x = 4
    .quad     .L7    # x = 5
    .quad     .L7    # x = 6
```

x86-64 Basics – Switch Statement

Jump table

```
.section .rodata
.align 8
.L4:
.quad .L8 # x = 0
.quad .L3 # x = 1
.quad .L5 # x = 2
.quad .L9 # x = 3
.quad .L8 # x = 4
.quad .L7 # x = 5
.quad .L7 # x = 6
```

```
switch(x) {
case 1:      // .L3
    w = y*z;
    break;
case 2:      // .L5
    w = y/z;
    /* Fall Through */
case 3:      // .L9
    w += z;
    break;
case 5:
case 6:      // .L7
    w -= z;
    break;
default:    // .L8
    w = 2;
}
```



x86-64 Basics – Switch Statement

Code Blocks (x == 2, x == 3)

```

long w = 1;
. . .
switch(x) {
. . .
case 2:
    w = y/z;
    /* Fall Through */
case 3:
    w += z;
    break;
. . .
}

```

```

.L5:                                # Case 2
    movq    %rsi, %rax
    cqto
    idivq   %rcx                    # y/z
    jmp     .L6                    # goto merge
.L9:                                # Case 3
    movl    $1, %eax               # w = 1
.L6:                                # merge:
    addq    %rcx, %rax             # w += z
    ret


```

Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rdx	Argument z
%rax	Return value

x86-64 Basics – Structs/Data align

Structure Representation

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- Structure represented as block of memory
 - Big enough to hold all of the fields
- Fields ordered according to declaration
 - Even if another ordering could yield a more compact representation
- Compiler determines overall size + positions of fields
 - Machine-level program has no understanding of the structures in the source code

x86-64 Basics – Structs/Data align

Alignment Principles

🌀 Aligned Data

- 🌀 Primitive data type requires K bytes
- 🌀 Address must be multiple of K
- 🌀 Required on some machines; advised on x86-64

🌀 Motivation for Aligning Data

- 🌀 Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - 🌀 Inefficient to load or store datum that spans quad word boundaries
 - 🌀 Virtual memory trickier when datum spans 2 pages

🌀 Compiler

- 🌀 Inserts gaps in structure to ensure correct alignment of fields

x86-64 Basics – Structs/Data align

Specific Cases of Alignment (x86-64)

- 1 byte: **char**, ...

- no restrictions on address

- 2 bytes: **short**, ...

- lowest 1 bit of address must be 0_2

- 4 bytes: **int**, **float**, ...

- lowest 2 bits of address must be 00_2

- 8 bytes: **double**, **long**, **char ***, ...

- lowest 3 bits of address must be 000_2

- 16 bytes: **long double** (GCC on Linux)

- lowest 4 bits of address must be 0000_2

x86-64 Basics – Structs/Data align

Satisfying Alignment with Structures

Within structure:

- Must satisfy each element's alignment requirement

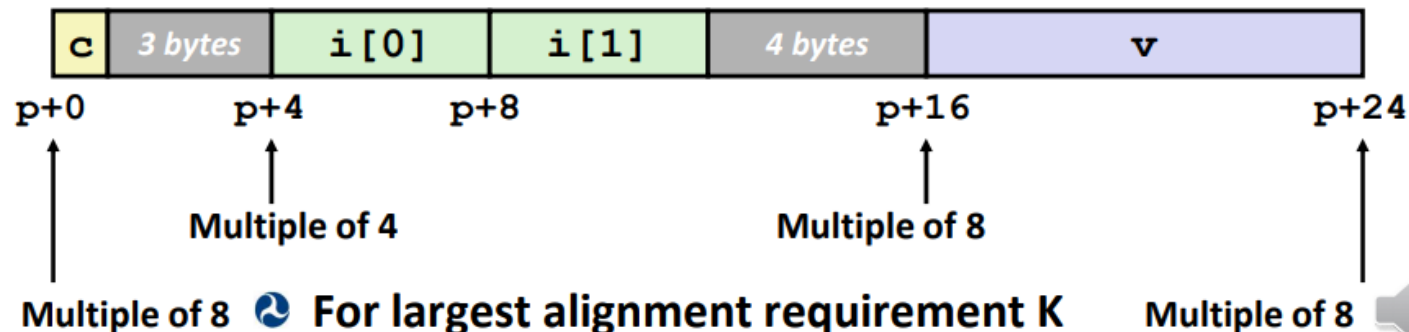
Overall structure placement

- Each structure has alignment requirement K
 - K = Largest alignment of any element
- Initial address & structure length must be multiples of K

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

Example:

- $K = 8$, due to `double` element



Multiple of 8 **For largest alignment requirement K**

Multiple of 8

- Overall structure must be multiple of K**

x86-64 Basics – Structs/Data align

Satisfying Alignment with Structures

Within structure:

- Must satisfy each element's alignment requirement

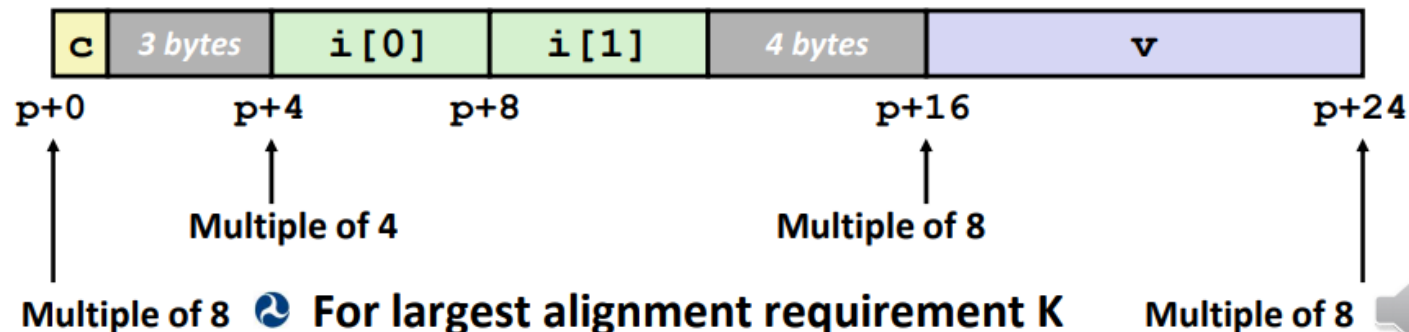
Overall structure placement

- Each structure has alignment requirement K
 - K = Largest alignment of any element
- Initial address & structure length must be multiples of K

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

Example:

- $K = 8$, due to `double` element



Multiple of 8

- For largest alignment requirement K

Multiple of 8

- Overall structure must be multiple of K

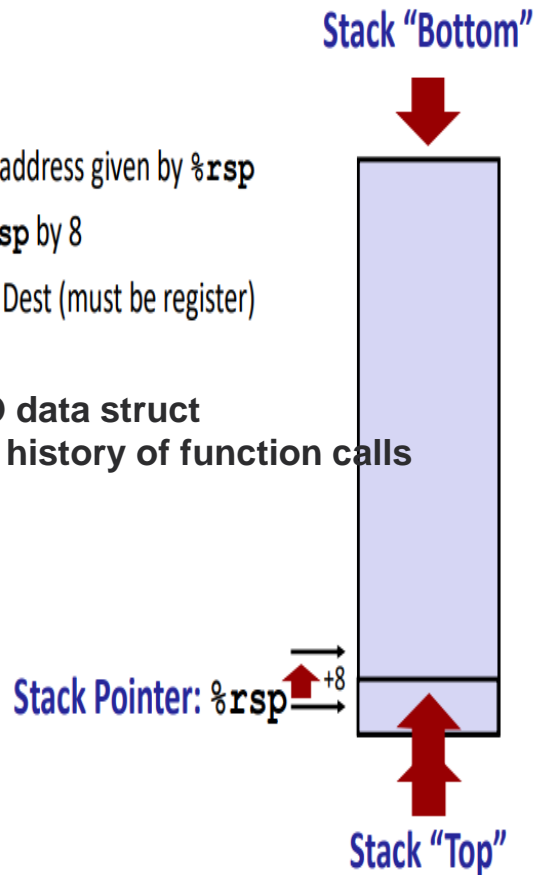
x86-64 Basics – Calling Convention

x86-64 Stack: Pop

■ `popq Dest`

- Read value at address given by `%rsp`
- Increment `%rsp` by 8
- Store value at `Dest` (must be register)

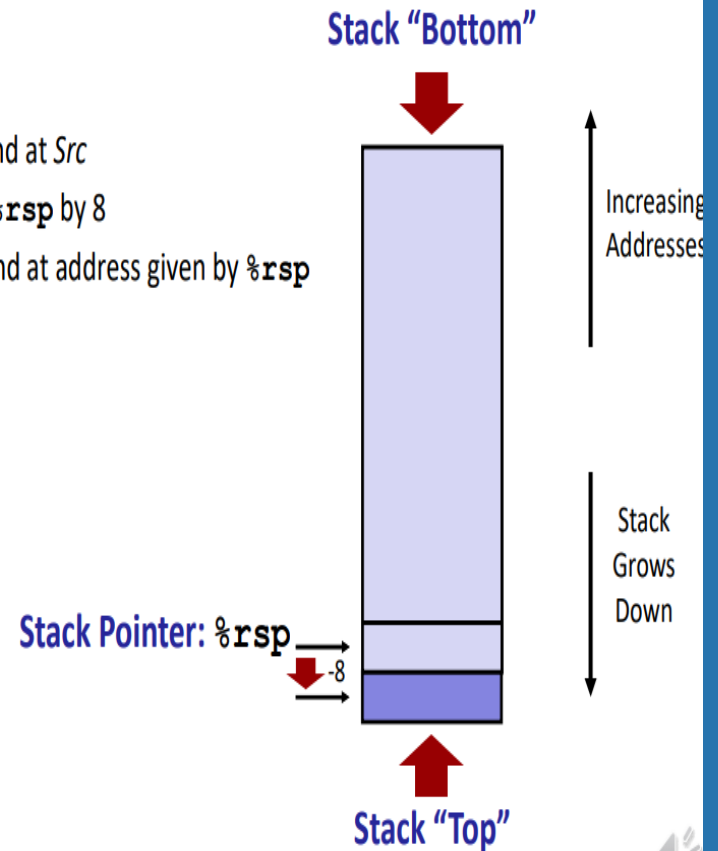
Stack is LIFO data struct
Contains the history of function calls



x86-64 Stack: Push

⌚ `pushq Src`

- ⌚ Fetch operand at `Src`
- ⌚ Decrement `%rsp` by 8
- ⌚ Write operand at address given by `%rsp`



x86-64 Basics – Calling Convention

Stack Frames

Contents


- Return information
- Local storage (if needed)
- Temporary space (if needed)

Management

- Space allocated when enter procedure
 - “Set-up” code
 - Includes push by **call** instruction
- Deallocated when return
 - “Finish” code
 - Includes pop by **ret** instruction

Frame Pointer: `%rbp`
(Optional)

Stack Pointer: `%rsp`




Stack “Top”



x86-64 Basics – Calling Convention

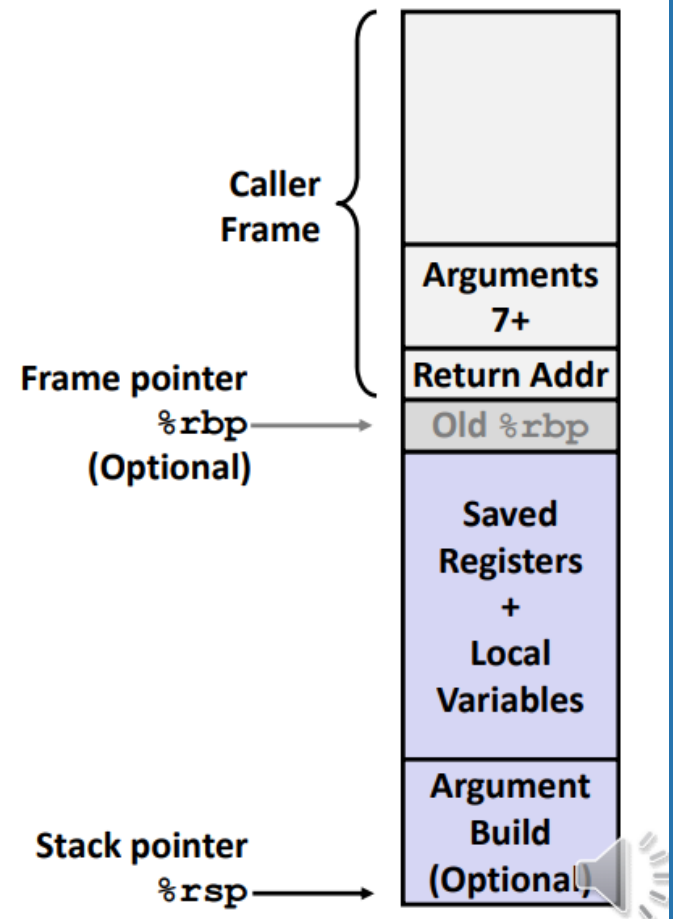
x86-64/Linux Stack Frame

Current Stack Frame (“Top” to Bottom)

- “Argument build:”
Parameters for function about to call
- Local variables
If can’t keep in registers
- Saved register context
- Old frame pointer (optional)

Caller Stack Frame

- Return address
 - Pushed by `call` instruction
- Arguments for this call



x86-64 Basics – Calling Convention

Register Saving Conventions

🌀 When procedure `yoo` calls `who`:

🌀 `yoo` is the *caller*

🌀 `who` is the *callee*

🌀 Can register be used for temporary storage?

```
yoo:
    . . .
    movq $15213, %rdx
    call who
    addq %rdx, %rax
    . . .
    ret
```

```
who:
    . . .
    subq $18213, %rdx
    . . .
    ret
```

🌀 Contents of register `%rdx` overwritten by `who`

🌀 This could be trouble → something should be done!

🌀 Need some coordination

x86-64 Basics – Calling Convention

Register Saving Conventions

🌀 When procedure `yoo` calls `who`:

🌀 `yoo` is the *caller*

🌀 `who` is the *callee*

🌀 Can register be used for temporary storage?

```
yoo:
    . . .
    movq $15213, %rdx
    call who
    addq %rdx, %rax
    . . .
    ret
```

```
who:
    . . .
    subq $18213, %rdx
    . . .
    ret
```

🌀 Contents of register `%rdx` overwritten by `who`

🌀 This could be trouble → something should be done!

🌀 Need some coordination

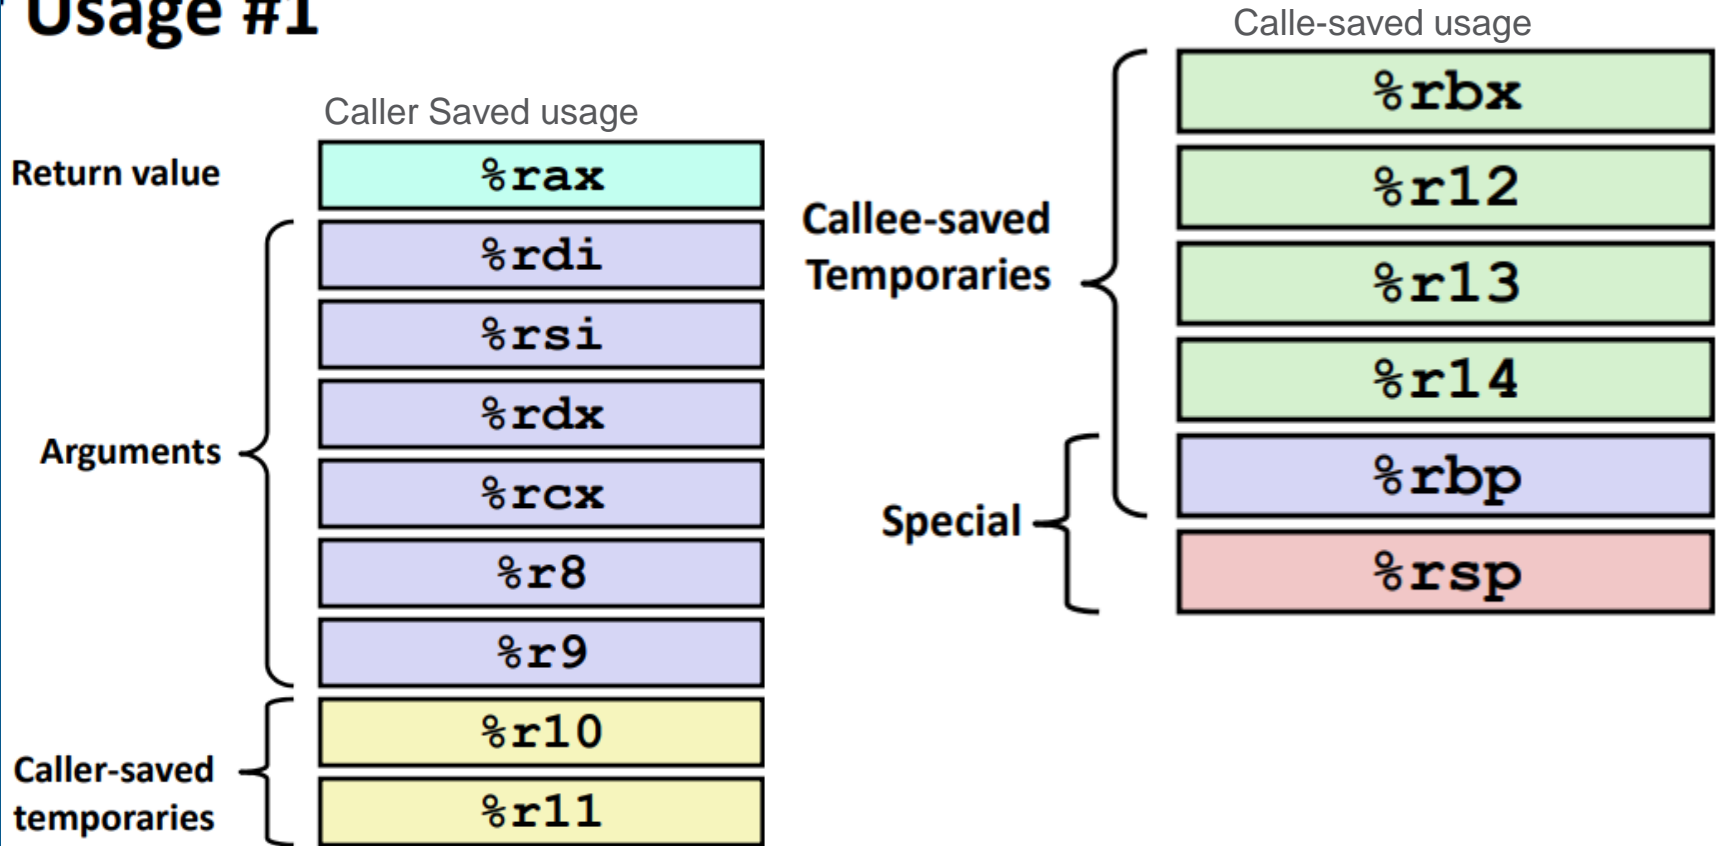
x86-64 Basics – Calling Convention

Register Saving Conventions

- When procedure `yoo` calls `who`:
 - `yoo` is the *caller*
 - `who` is the *callee*
- Can register be used for temporary storage?
- Conventions
 - “Caller Saved”*
 - Caller saves temporary values in its frame before the call
 - “Callee Saved”*
 - Callee saves temporary values in its frame before using
 - Callee restores them before returning to caller

x86-64 Basics – Calling Convention

Usage #1




x86-64 Basics – Recursion

Recursive Function Register Save

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

```
pcount_r:
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq    %rdi, %rbx
    andl    $1, %ebx
    shrq    %rdi
    call    pcount_r
    addq    %rbx, %rax
    popq    %rbx
.L6:
    ret
```

Register	Use(s)	Type
%rdi	x	Argument



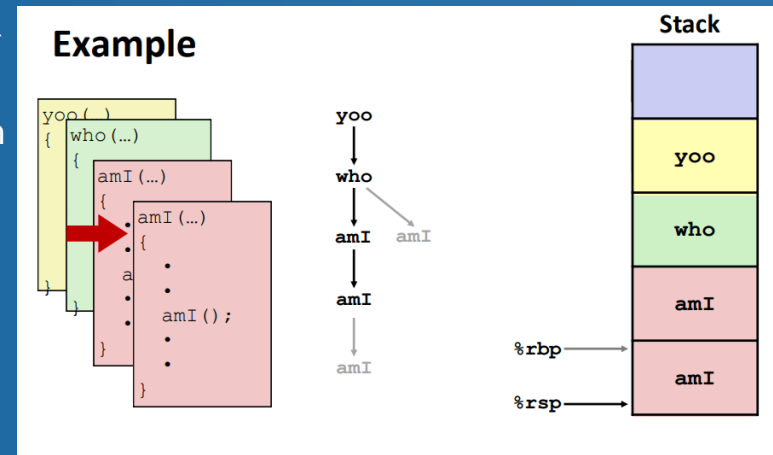
Notes adapted from Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

48

x86-64 Basics – Recursion

- A potential problem: stack overflow
- To reduce the chance of stack overflow:
- Tail-recursion: make the last instruction in the function be the recursive call (last instruction is the function calling itself)
- Tail recursive function calls don't need to return to caller, since there's nothing left in the caller to be handled.
 - No need to store the caller's address
 - No need to store different stack frames for each invocation
 - Saves tons of space on stack

Example of stack with regular function calls, can easily overload with stack frames :



x86-64 Basics – Recursion

- How do we make this function tail recursive?
- What does the stack look like in regular / tail recursive example?
- Lets investigate!
- Compiler should optimize recursion into loop

```

00000000004004c4 <factorial>:
4004c4: 53                push    %rbx
4004c5: 89 fb            mov     %edi,%ebx
4004c7: b8 00 00 00 00   mov     $0x0,%eax
4004cc: 85 ff            test    %edi,%edi
4004ce: 74 12            je      4004e2 <factorial+0x1e>
4004d0: b0 01            mov     $0x1,%al
4004d2: 83 ff 01         cmp     $0x1,%edi
4004d5: 74 0b            je      4004e2 <factorial+0x1e>
4004d7: 8d 7b ff         lea     -0x1(%rbx),%edi
4004da: e8 e5 ff ff ff   callq   4004c4 <factorial>
4004df: 0f af c3         imul    %ebx,%eax
4004e2: 5b              pop     %rbx
4004e3: c3              retq

00000000004004e4 <main>:
4004e4: 48 83 ec 08      sub     $0x8,%rsp
4004e8: bf 03 00 00 00   mov     $0x3,%edi
4004ed: e8 d2 ff ff ff   callq   4004c4 <factorial>
4004f2: 89 c6            mov     %eax,%esi
4004f4: bf 08 06 40 00   mov     $0x400608,%edi
4004f9: b8 00 00 00 00   mov     $0x0,%eax
4004fe: e8 b5 fe ff ff   callq   4003b8 <printf@plt>
400503: b8 01 00 00 00   mov     $0x1,%eax
400508: 48 83 c4 08      add     $0x8,%rsp
40050c: c3              retq
40050d: 90              nop
40050e: 90              nop
40050f: 90              nop

```

```

#include <stdlib.h>
#include <stdio.h>

unsigned int factorial (unsigned int x)
{
    if (x==0)
        return 0;
    else
        if (x==1)
            return 1;
        else
            return x*factorial(x-1);
}

int main( int argc, const char* argv[] )
{
    printf("%d\n", factorial(3));

    return 1;
}

```


Questions before we start worksheet?

Resources used

- https://cs.brown.edu/courses/cs033/docs/guides/x64_cheatsheet.pdf
- <https://cs.stackexchange.com/questions/76871/how-are-variables-stored-in-and-retrieved-from-the-program-stack>
- <https://www.geeksforgeeks.org/memory-layout-of-c-program/>
- <https://vim.rtorr.com/>
- <https://www.tenouk.com/ModuleW.html>
- Professor Reinman's slides (CCLE)