

CS 33 – Discussion 1A

Computer System Organization

Week 5

Questions before we start?

Logistics

- Lab 1 - Bomblab due tonight 11:59pm
- Lab 2 – Attack lab will be released May 5th – Due Friday May 18th 11:55pm
- HW 3 will be released soon, due May 10th
- Please submit suggestions on CCLE TA-site on how to improve OH or DISC
 - Please fill out LA survey on CCLE

Agenda

10am-11am PST:

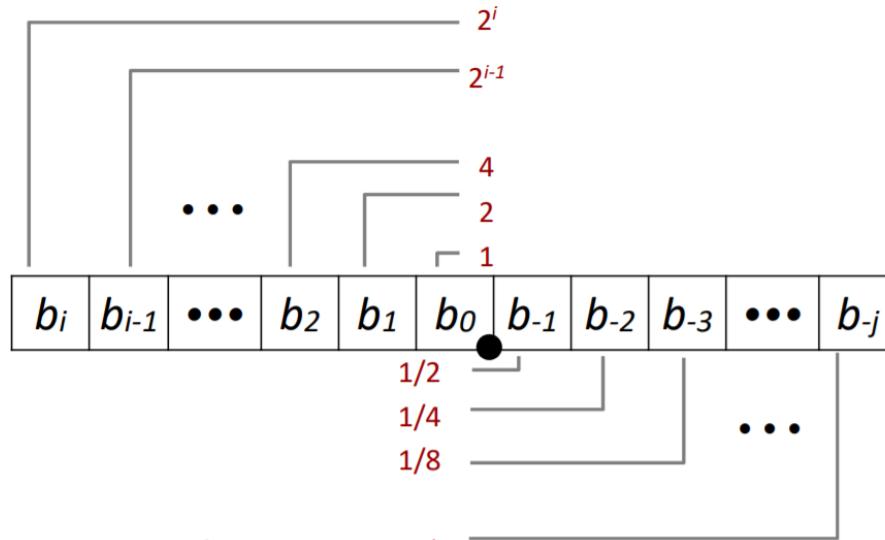
- Go through midterm together?
- Floating point
- Stack overflow attacks
- Optimization

11am-11:50am PST:

- LA worksheet

x86-64 Basics – Floating point

Fractional Binary Numbers



Representation

- Bits to right of “binary point” represent fractional powers of 2
- Represents rational number:
$$\sum_{k=-j}^i b_k \times 2^k$$

x86-64 Basics – Floating point

Floating Point Representation

➊ Numerical Form:

$$(-1)^s M \ 2^E$$

- ➊ **Sign bit *s*** determines whether number is negative or positive
- ➋ **Significand *M*** normally a fractional value in range [1.0,2.0).
- ➌ **Exponent *E*** weights value by power of two

➋ Encoding

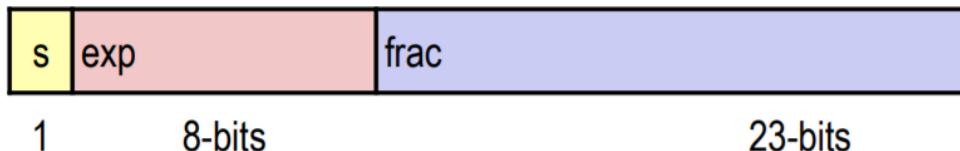
- ➊ MSB *s* is sign bit ***s***
- ➋ exp field encodes ***E*** (but is not equal to *E*)
- ➌ frac field encodes ***M*** (but is not equal to *M*)



x86-64 Basics – Floating point

Precision options

➊ Single precision: 32 bits



➋ Double precision: 64 bits



➌ Extended precision: 80 bits (Intel only)



➍ Conversions/Casting

- ➎ Casting between `int`, `float`, and `double` changes bit representation
- ➏ `double/float → int`
 - ➐ Truncates fractional part
 - ➑ Like rounding toward zero
 - ➒ Not defined when out of range or NaN: Generally sets to TMin
- ➏ `int → double`
 - ➐ Exact conversion, as long as `int` has \leq 53 bit word size
- ➏ `int → float`
 - ➐ Will round according to rounding mode

x86-64 Basics – Floating point

“Normalized” Values

$$v = (-1)^s M 2^E$$

- ➊ When: $\text{exp} \neq 000\ldots0$ and $\text{exp} \neq 111\ldots1$
- ➋ Exponent coded as a *biased value*: $E = \text{Exp} - \text{Bias}$
 - ➌ Exp : unsigned value of exp field
 - ➌ $\text{Bias} = 2^{k-1} - 1$, where k is number of exponent bits
 - ➌ Single precision: 127 ($\text{Exp}: 1\ldots254$, $E: -126\ldots127$)
 - ➌ Double precision: 1023 ($\text{Exp}: 1\ldots2046$, $E: -1022\ldots1023$)
- ➌ Significand coded with implied leading 1: $M = 1.\text{xxx}\ldots\text{x}_2$
 - ➌ $\text{xxx}\ldots\text{x}$: bits of frac field
 - ➌ Minimum when $\text{frac}=000\ldots0$ ($M = 1.0$)
 - ➌ Maximum when $\text{frac}=111\ldots1$ ($M = 2.0 - \epsilon$)
 - ➌ Get extra leading bit for “free”



x86-64 Basics – Floating point

Denormalized Values

$$v = (-1)^s M 2^E$$
$$E = 1 - \text{Bias}$$

- ⌚ Condition: $\text{exp} = 000\dots0$
- ⌚ Exponent value: $E = 1 - \text{Bias}$ (instead of $E = 0 - \text{Bias}$)
- ⌚ Significand coded with implied leading 0: $M = 0.\text{xxx}\dots\text{x}_2$
 - ⌚ $\text{xxx}\dots\text{x}$: bits of `frac`
- ⌚ Cases
 - ⌚ $\text{exp} = 000\dots0, \text{frac} = 000\dots0$
 - ⌚ Represents zero value
 - ⌚ Note distinct values: +0 and -0
 - ⌚ $\text{exp} = 000\dots0, \text{frac} \neq 000\dots0$
 - ⌚ Numbers closest to 0.0
 - ⌚ Equispaced

x86-64 Basics – Floating point

Special Values

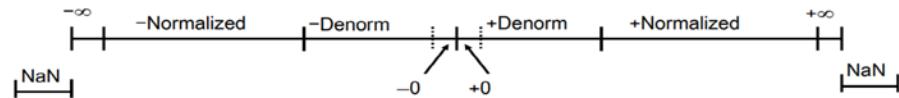
- ➊ Condition: **exp = 111...1**

- ➋ Case: **exp = 111...1, frac = 000...0**
 - ➌ Represents value ∞ (infinity)
 - ➌ Operation that overflows
 - ➌ Both positive and negative
 - ➌ E.g., $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$

- ➌ Case: **exp = 111...1, frac \neq 000...0**
 - ➌ Not-a-Number (NaN)
 - ➌ Represents case when no numeric value can be determined
 - ➌ E.g., $\sqrt{-1}$, $\infty - \infty$, $\infty \times 0$

x86-64 Basics – Floating point

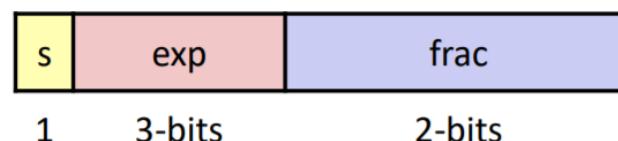
Visualization: Floating Point Encodings



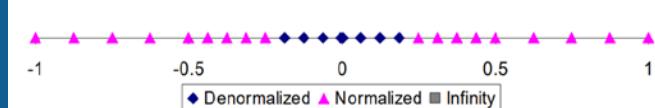
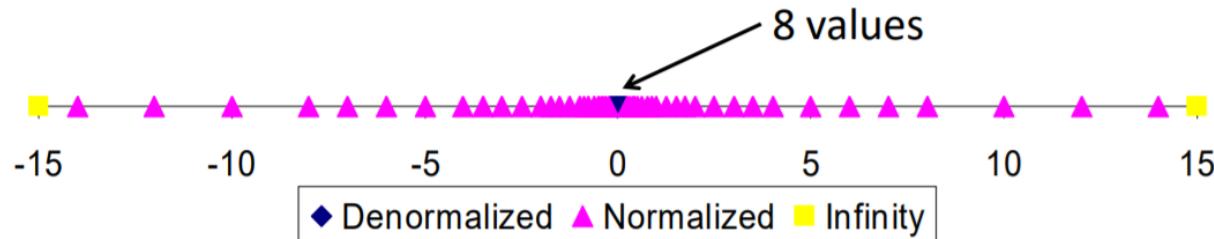
Distribution of Values

6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- Bias is $2^{3-1}-1 = 3$



Notice how the distribution gets denser toward zero.



x86-64 Basics – Floating point

Special Properties of the IEEE Encoding

⌚ FP Zero Same as Integer Zero

- ⌚ All bits = 0

⌚ Can (Almost) Use Unsigned Integer Comparison

- ⌚ Must first compare sign bits
- ⌚ Must consider $-0 = 0$
- ⌚ NaNs problematic
 - ⌚ Will be greater than any other values
 - ⌚ What should comparison yield?
- ⌚ Otherwise OK
 - ⌚ Denorm vs. normalized
 - ⌚ Normalized vs. infinity

Floating Point Operations: Basic Idea

$$\⌚ x +_f y = \text{Round}(x + y)$$

$$\⌚ x \times_f y = \text{Round}(x \times y)$$

⌚ Basic idea

- ⌚ First **compute exact result**
- ⌚ Make it fit into desired precision
 - ⌚ Possibly overflow if exponent too large
 - ⌚ Possibly **round to fit into `frac`**

x86-64 Basics – Floating point

Rounding Binary Numbers

Binary Fractional Numbers

- “Even” when least significant bit is 0
- “Half way” when bits to right of rounding position = $100\dots_2$

Examples

- Round to nearest 1/4 (2 bits right of binary point)

Value Value	Binary	Rounded	Action	Rounded
2 3/32	10.00011 ₂	10.00 ₂	(<1/2—down)	2
2 3/16	10.00110 ₂	10.01 ₂	(>1/2—up)	2 1/4
2 7/8	10.11100 ₂	11.00 ₂	(1/2—up)	3
2 5/8	10.10100 ₂	10.10 ₂	(1/2—down)	2 1/2

x86-64 Basics – Floating point

FP Multiplication

⌚ $(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$

⌚ **Exact Result:** $(-1)^s M 2^E$

⌚ Sign s : $s1 \wedge s2$

⌚ Significand M : $M1 \times M2$

⌚ Exponent E : $E1 + E2$

⌚ Fixing

⌚ If $M \geq 2$, shift M right, increment E

⌚ If E out of range, overflow

⌚ Round M to fit `frac` precision

⌚ Implementation

⌚ Biggest chore is multiplying significands

x86-64 Basics – Floating point

Floating Point Addition

➊ $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$

➋ Assume $E1 > E2$

➌ Exact Result: $(-1)^s M 2^E$

➍ Sign s , significand M :

➎ Result of signed align & add

➏ Exponent E : $E1$

➐ Fixing

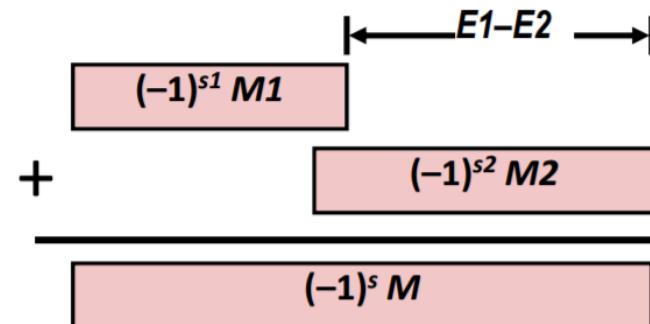
➑ If $M \geq 2$, shift M right, increment E

➒ if $M < 1$, shift M left k positions, decrement E by k

➓ Overflow if E out of range

➔ Round M to fit `frac` precision

Get binary points lined up



x86-64 Basics – Floating point

Floating Point Addition

➊ $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$

➋ Assume $E1 > E2$

➌ Exact Result: $(-1)^s M 2^E$

➍ Sign s , significand M :

➎ Result of signed align & add

➏ Exponent E : $E1$

➐ Fixing

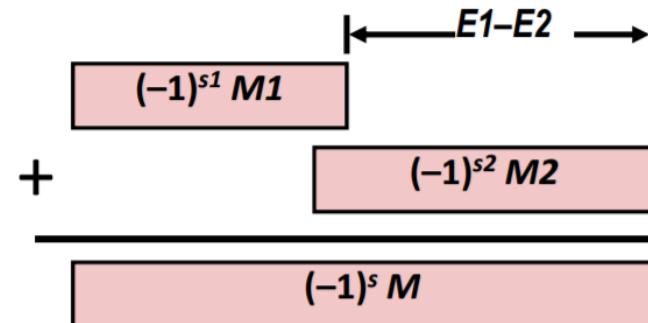
➑ If $M \geq 2$, shift M right, increment E

➒ if $M < 1$, shift M left k positions, decrement E by k

➓ Overflow if E out of range

➔ Round M to fit `frac` precision

Get binary points lined up



x86-64 Basics – Stack attacks

- Stack overflow attacks
- Inject code on stack to execute on stack
 - Through buffer overflow into local variables
 - e.g. int buffer[LIMIT] with |input| > LIMIT
 - Stick byte-code that corresponds to your ins
 - When it returns from function, pops ret val
 - Ret val overwritten to beginning of stack
- Return oriented programming
 - Find desired byte code sequences for ins in text memory
 - Must end with c3 – for retq
 - Stick pointer to code on stack beginning at overwritten ret addr (address)

x86-64 Basics – Stack attacks

- Stack overflow attack protections
- For first type:
 - make stack non-executeable
 - Permission bits on pages, can be hardware check
 - Randomize stack location
- For ROP: use stack canaries
 - Special hidden value placed on stack after ret addr
 - See if ret addr overwritten with canary check
 - Then return from function else raise fault

x86-64 Basics – Stack attacks

Recall: Memory Referencing Bug Example

Explanation:

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

Critical State	6
?	5
?	4
d7 ... d4	3
d3 ... d0	2
a[1]	1
a[0]	0

```
fun(0)  => 3.14
fun(1)  => 3.14
fun(2)  => 3.1399998664856
fun(3)  => 2.00000061035156
fun(4)  => 3.14
fun(6)  => Segmentation fault
```

x86-64 Basics – Stack attacks

Vulnerable Buffer Code

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

←btw, how big
is big enough?

```
void call_echo() {
    echo();
}
```

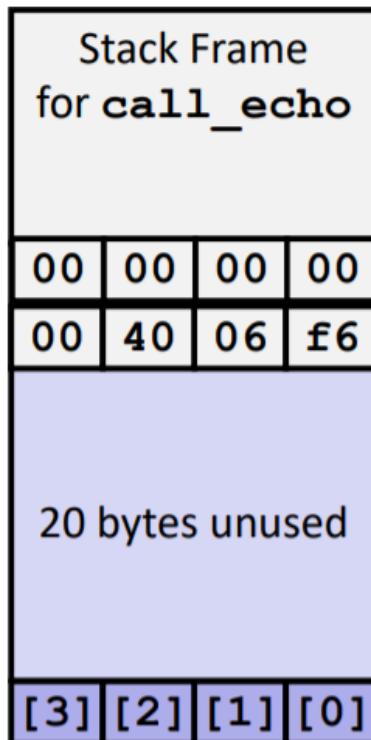
```
unix>./bufdemo-nsp
Type a string:012345678901234567890123
012345678901234567890123
```

```
unix>./bufdemo-nsp
Type a string:0123456789012345678901234
Segmentation Fault
```

x86-64 Basics – Stack attacks

Buffer Overflow Stack Example

Before call to gets



```
void echo()
{
    char buf[4];
    gets(buf);
    ...
}
```

```
echo:
    subq $24, %rsp
    movq %rsp, %rdi
    call gets
    ...
}
```

`call_echo:`

```
...
4006f1: callq 4006cf <echo>
4006f6: add    $0x8,%rsp
...
```

x86-64 Basics – Stack attacks

```
unix>./bufdemo-nsp  
Type a string:01234567890123456789012  
01234567890123456789012
```

After call to gets

Stack Frame
for `call_echo`

00	00	00	00
00	40	06	f6
00	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

Could overwrite
stack to jump to
address that doesn't
cause seg fault

```
unix>./bufdemo-nsp  
Type a string:0123456789012345678901234  
Segmentation Fault
```

After call to gets

Stack Frame
for `call_echo`

00	00	00	00
00	40	00	34
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

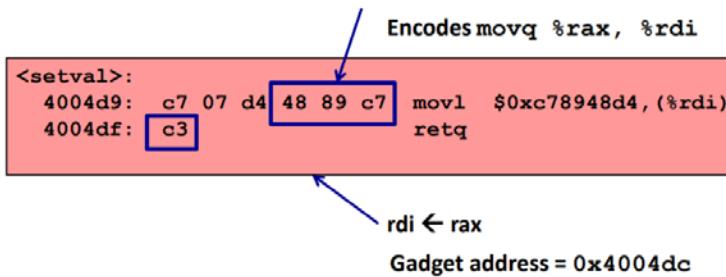
x86-64 Basics – Stack attacks ROP

Construct program from *gadgets*

- Sequence of instructions ending in `ret`
 - Encoded by single byte `0xc3`
- Code positions fixed from run to run
- Code is executable

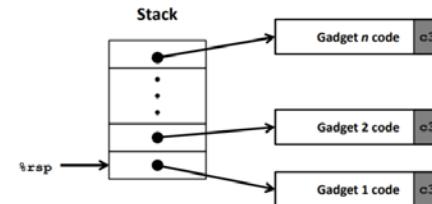
Gadget Example #2

```
void setval(unsigned *p) {
    *p = 3347663060u;
}
```



- Repurpose byte codes

ROP Execution



- Trigger with `ret` instruction
 - Will start executing Gadget 1
- Final `ret` in each gadget will start next one

Gadget Example #1

```
long ab_plus_c
(long a, long b, long c)
{
    return a*b + c;
}
```

```

000000000004004d0 <ab_plus_c>:
4004d0: 48 0f af fe    imul %rsi,%rdi
4004d4: 48 8d 04 17    lea (%rdi,%rdx,1),%rax
4004d8: c3              retq

```

`rax \leftarrow rdi + rdx`
Gadget address = `0x4004d4`

- Use tail end of existing functions

x86-64 Basics – Optimization

Basic types of optimization:

- Code motion
- Reduction in Strength
- Share common Subexpressions
- Loop unrolling

Optimization blockers:

- Procedure calls
 - Fix by doing own code motion
 - Inline functions
- Memory Aliasing
 - Pointers can point to same object
 - Makes it so you have to access memory everytime to mutate object
 - Keep a local copy to amortize reads/writes
 - Accumulate in loop

x86-64 Basics – Code motion

■ Code Motion

- Reduce frequency with which computation performed
 - If it will always produce same result
 - Especially moving code out of loop

```
void set_row(double *a, double *b,
             long i, long n)
{
    long j;
    for (j = 0; j < n; j++)
        a[n*i+j] = b[j];
}
```



```
long j;
int ni = n*i;
for (j = 0; j < n; j++)
    a[ni+j] = b[j];
```

x86-64 Basics – Code motion

■ Code Motion

- Reduce frequency with which computation performed
 - If it will always produce same result
 - Especially moving code out of loop

```
void set_row(double *a, double *b,
            long i, long n)
{
    long j;
    for (j = 0; j < n; j++)
        a[n*i+j] = b[j];
}
```

```
long j;
int ni = n*i;
for (j = 0; j < n; j++)
    a[ni+j] = b[j];
```

Compiler-Generated Code Motion (-O1)

```
void set_row(double *a, double *b,
            long i, long n)
{
    long j;
    for (j = 0; j < n; j++)
        a[n*i+j] = b[j];
}
```

```
long j;
long ni = n*i;
double *rowp = a+ni;
for (j = 0; j < n; j++)
    *rowp++ = b[j];
```

```
set_row:
    testq    %rcx, %rcx          # Test n
    jle     .L1                  # If 0, goto done
    imulq   %rcx, %rdx          # ni = n*i
    leaq    (%rdi,%rdx,8), %rdx # rowp = A + ni*8
    movl    $0, %eax             # j = 0
    .L3:
    movsd   (%rsi,%rax,8), %xmm0 # t = b[j]
    movsd   %xmm0, (%rdx,%rax,8) # M[A+ni*8 + j*8] = t
    addq    $1, %rax              # j++
    cmpq    %rcx, %rax          # j:n
    jne     .L3                  # if !=, goto loop
    .L1:
    rep : ret
```

x86-64 Basics – share calculation

Share Common Subexpressions

- Reuse portions of expressions
- GOC will do this with -O1

```
/* Sum neighbors of i,j */
up = val[(i-1)*n + j];
down = val[(i+1)*n + j];
left = val[i*n      + j-1];
right = val[i*n      + j+1];
sum = up + down + left + right;
```

3 multiplications: $i*n$, $(i-1)*n$, $(i+1)*n$

```
leaq  1(%rsi), %rax # i+1
leaq -1(%rsi), %r8  # i-1
imulq %rcx, %rsi    # i*n
imulq %rcx, %rax    # (i+1)*n
imulq %rcx, %r8     # (i-1)*n
addq  %rdx, %rsi    # i*n+j
addq  %rdx, %rax    # (i+1)*n+j
addq  %rdx, %r8     # (i-1)*n+j
```

```
long inj = i*n + j;
up = val[inj - n];
down = val[inj + n];
left = val[inj - 1];
right = val[inj + 1];
sum = up + down + left + right;
```

1 multiplication: $i*n$

```
imulq  %rcx, %rsi # i*n
addq  %rdx, %rsi # i*n+j
movq  %rsi, %rax # i*n+j
subq  %rcx, %rax # i*n+j-n
leaq  (%rsi,%rcx), %rcx # i*n+j+n
```

x86-64 Basics – Loop unrolling

Loop Unrolling (2x1)

```
void unroll2a_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x = (x OP d[i]) OP d[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x = x OP d[i];
    }
    *dest = x;
}
```

- Perform 2x more useful work per iteration

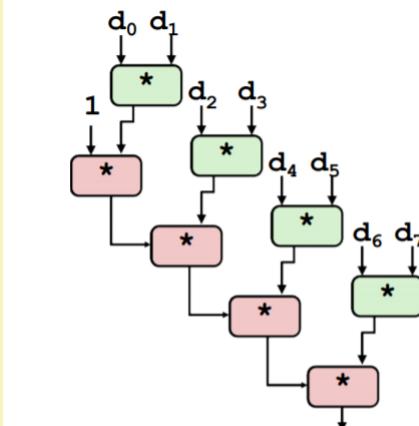
x86-64 Basics – Loop unrolling

Loop Unrolling with Reassociation (2x1a)

```
void unroll2aa_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x = x OP (d[i] OP d[i+1]);
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x = x OP d[i];
    }
    *dest = x;
}
```

Compare to before

```
x = (x OP d[i]) OP d[i+1];
```

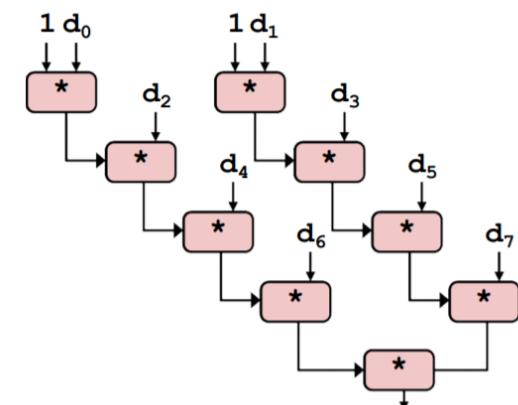


- Can this change the result of the computation?

x86-64 Basics – Loop unrolling

Loop Unrolling with Separate Accumulators (2x2)

```
void unroll2a_combine(vec_ptr v, data_t *dest)
{
    long length = vec_length(v);
    long limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x0 = IDENT;
    data_t x1 = IDENT;
    long i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x0 = x0 OP d[i];
        x1 = x1 OP d[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x0 = x0 OP d[i];
    }
    *dest = x0 OP x1;
}
```



x86-64 Basics – Optimization block

Optimization Blocker: Procedure Calls

■ *Why couldn't compiler move strlen out of inner loop?*

- Procedure may have side effects
 - Alters global state each time called
- Function may not return same value for given arguments
 - Depends on other parts of global state
 - Procedure `lower` could interact with `strlen`

■ **Warning:**

- Compiler treats procedure call as a black box
- Weak optimizations near them

■ **Remedies:**

- Use of inline functions
 - GCC does this with `-O1`
 - Within single file
- Do your own code motion

```
size_t lencnt = 0;
size_t strlen(const char *s)
{
    size_t length = 0;
    while (*s != '\0') {
        s++; length++;
    }
    lencnt += length;
    return length;
}
```

x86-64 Basics – Optimization block

Memory Aliasing

```
/* Sum rows is of n X n matrix a
   and store in vector b */
void sum_rows1(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        b[i] = 0;
        for (j = 0; j < n; j++)
            b[i] += a[i*n + j];
    }
}
```

FIX:

```
for (i = 0; i < n; i++) {
    double val = 0;
    for (j = 0; j < n; j++)
        val += a[i*n + j];
    b[i] = val;
}
```

```
double A[9] =
{ 0, 1, 2,
  4, 8, 16},
  32, 64, 128};

double B[3] = A+3;

sum_rows1(A, B, 3);
```

- Code updates `b[i]` on every iteration
- Must consider possibility that these updates will affect program behavior

■ Aliasing

- Two different memory references specify single location
- Easy to happen in C
 - Since allowed to do address arithmetic
 - Direct access to storage structures
- Get in habit of introducing local variables
 - Accumulating within loops
 - Your way of telling compiler not to check for aliasing

Value of B:

init:	[4, 8, 16]
i = 0:	[3, 8, 16]
i = 1:	[3, 22, 16]
i = 2:	[3, 22, 224]

Questions before we start worksheet?

Resources used

- Professor Reinman's slides (CCLE)
- **Credit for compilation: Attiano Purpura-Pontoniere**