

Agenda

- Virtual memory

Ideal Memory

- Zero access time (latency)
- **Infinite capacity**
- Zero cost
- Infinite bandwidth (to support multiple accesses in parallel)

Abstraction: Virtual vs. Physical Memory

- **Programmer** sees **virtual memory**
 - Can assume the memory is “infinite”
 - Reality: Physical memory size is much smaller than what the programmer assumes
 - **The system** (system software + hardware, cooperatively) **maps virtual memory addresses to physical memory**
 - The system automatically manages the physical memory space **transparently to the programmer**
- + Programmer does not need to know the physical size of memory nor manage it → A small physical memory can appear as a huge one to the programmer → Life is easier for the programmer
- More complex system software and architecture


Benefits of Automatic Management of Memory

- Programmer does not deal with physical addresses
- Each process has its own mapping from virtual → physical addresses
- Enables
 - Code and data to be located anywhere in physical memory
(relocation)
 - Isolation/separation of code and data of different processes in physical processes
(protection and isolation)
 - Code and data sharing between multiple processes
(sharing)

Basic Mechanism

- Indirection (in addressing)
- Address generated by each instruction in a program is a “virtual address”
 - i.e., it is not the physical address used to address main memory
- An “address translation” mechanism maps this address to a “physical address”
 - Address translation mechanism can be implemented in hardware and software together

A System with Virtual Memory (Page based)



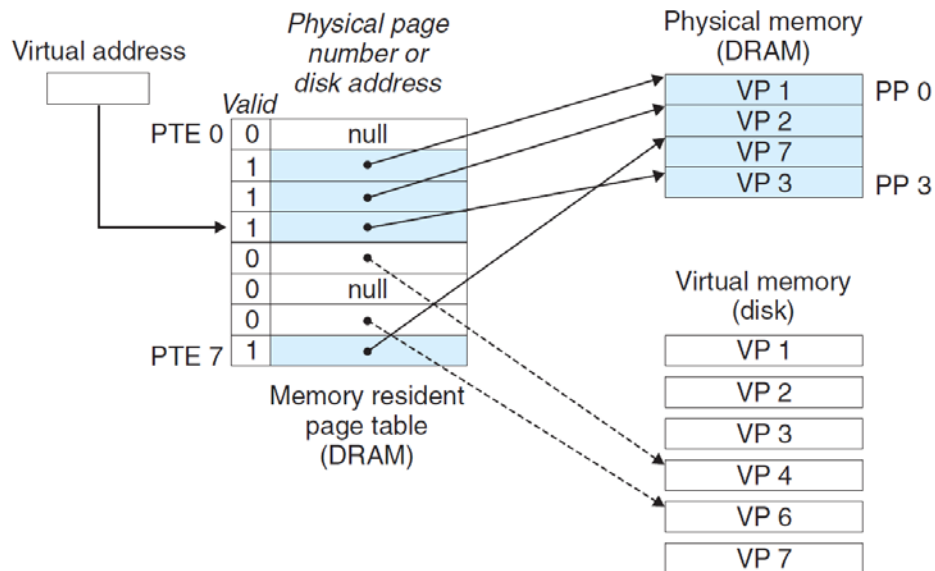
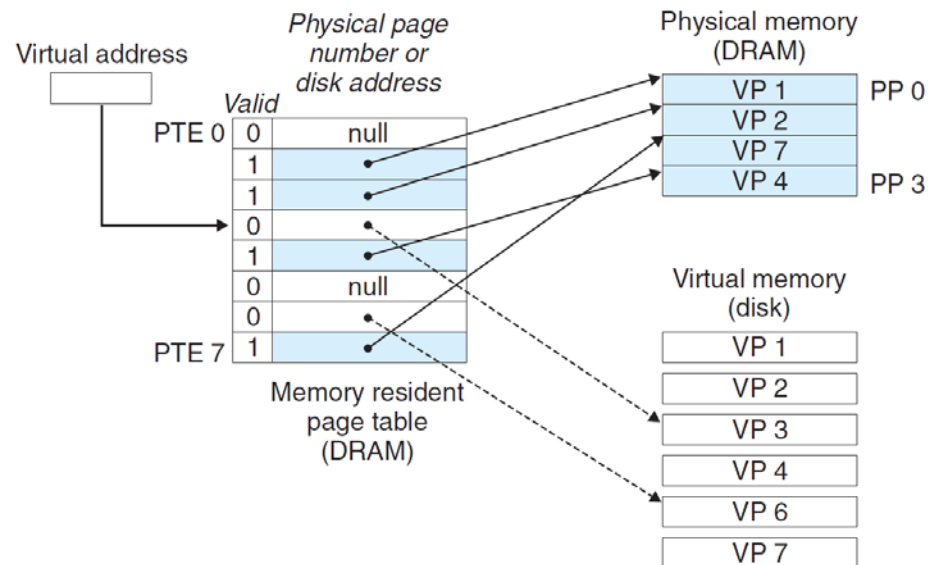
- Address Translation: The hardware converts virtual addresses into physical addresses via an OS-managed lookup table (page table)

Virtual Pages, Physical Pages

- Virtual and physical address space divided into **pages**
- A virtual page is mapped to
 - A physical page, if the page is in physical memory
 - A location in disk, otherwise
- If an accessed virtual page is not in memory, but on disk
 - Virtual memory system brings the page into a physical page and adjusts the mapping → this is called **demand paging**
- **Page table** is the table that stores the mapping of virtual pages to physical frames

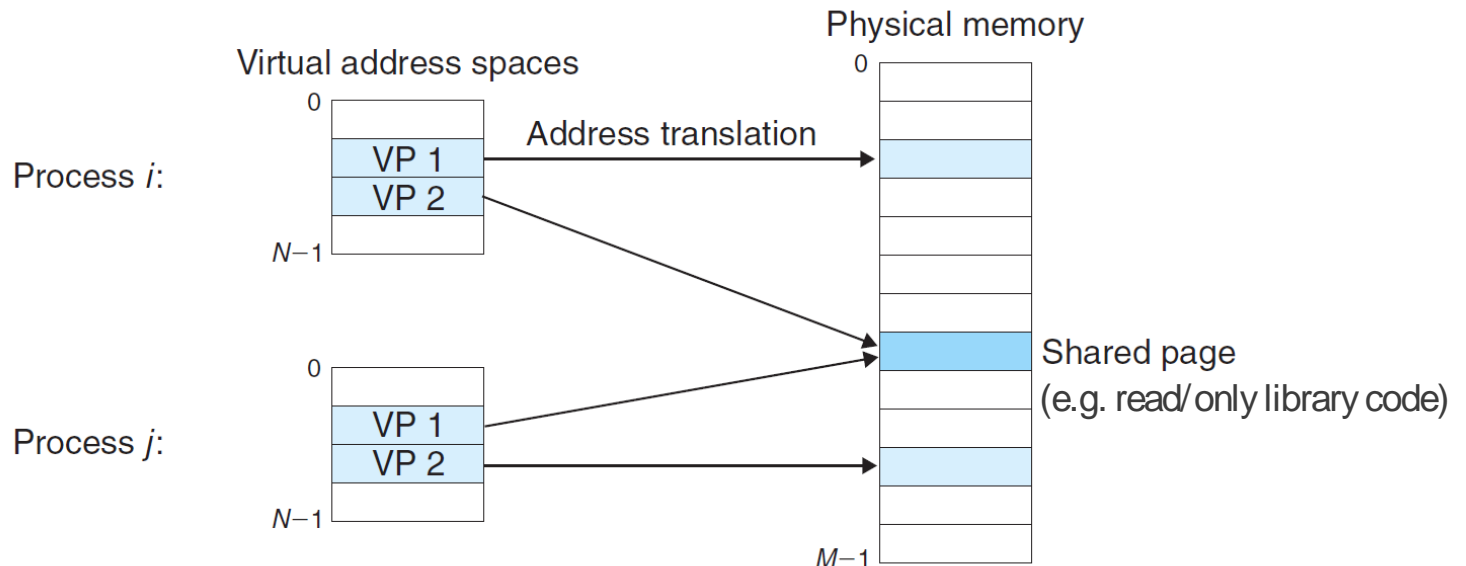
Page Fault (“A Miss in Physical Memory”)

- If a page is not in physical memory but disk
 - Page table entry indicates virtual page not in memory
 - Access to such a page triggers a page fault exception
 - OS trap handler invoked to move data from disk into memory
 - Other processes can continue executing
 - OS has full control over placement



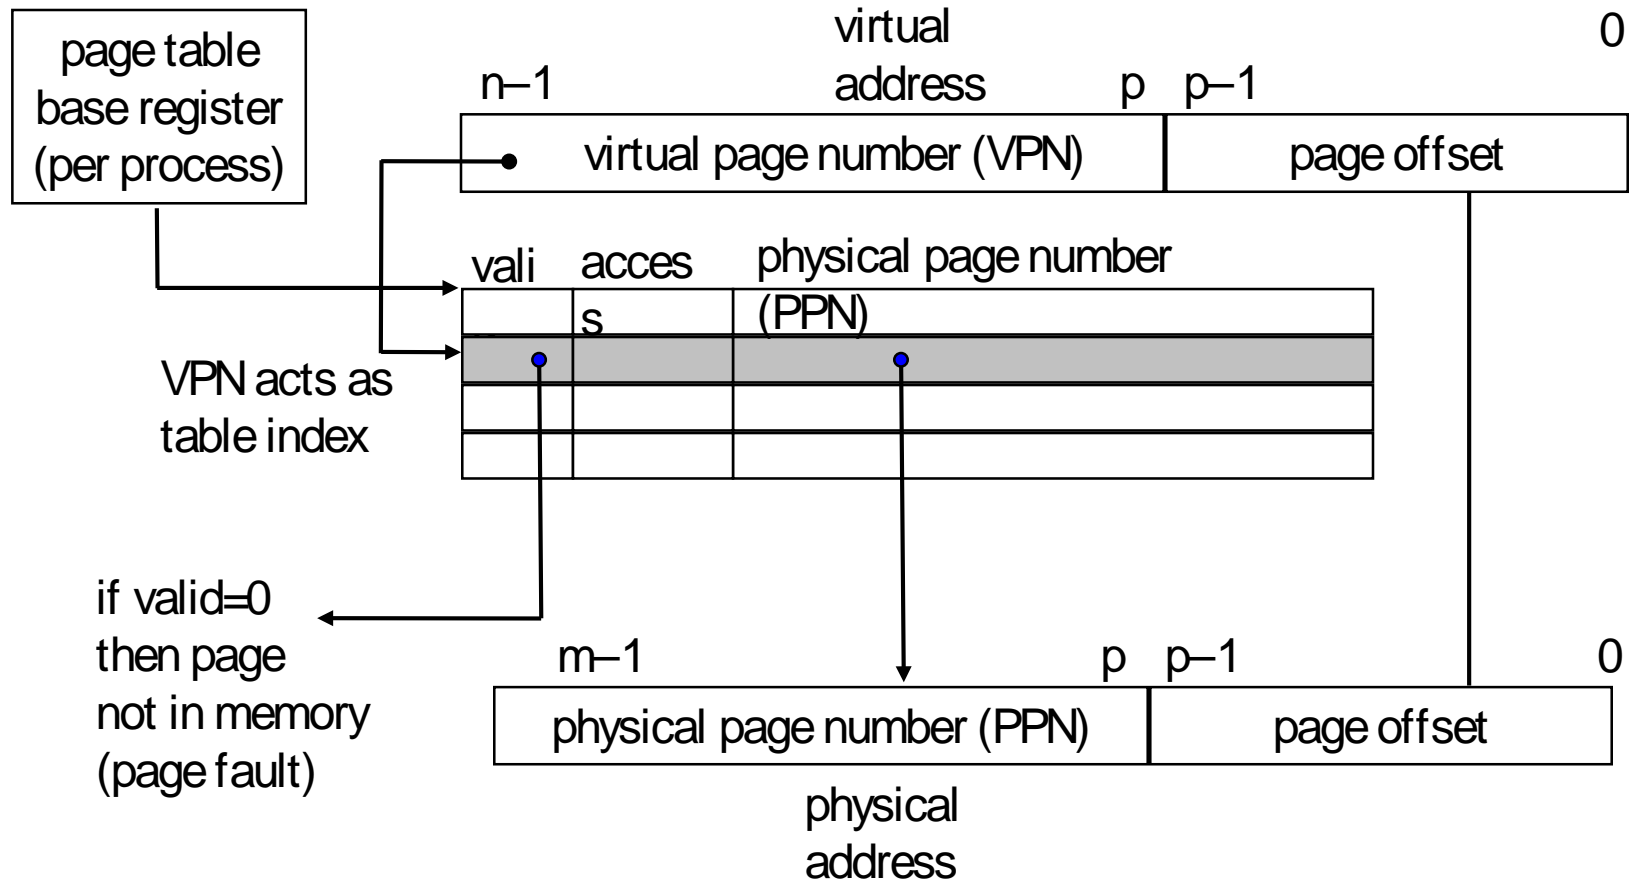
Page Table is Per Process

- Each process has its own virtual address space
 - Full address space for each program
 - Simplifies memory allocation, sharing, linking and loading



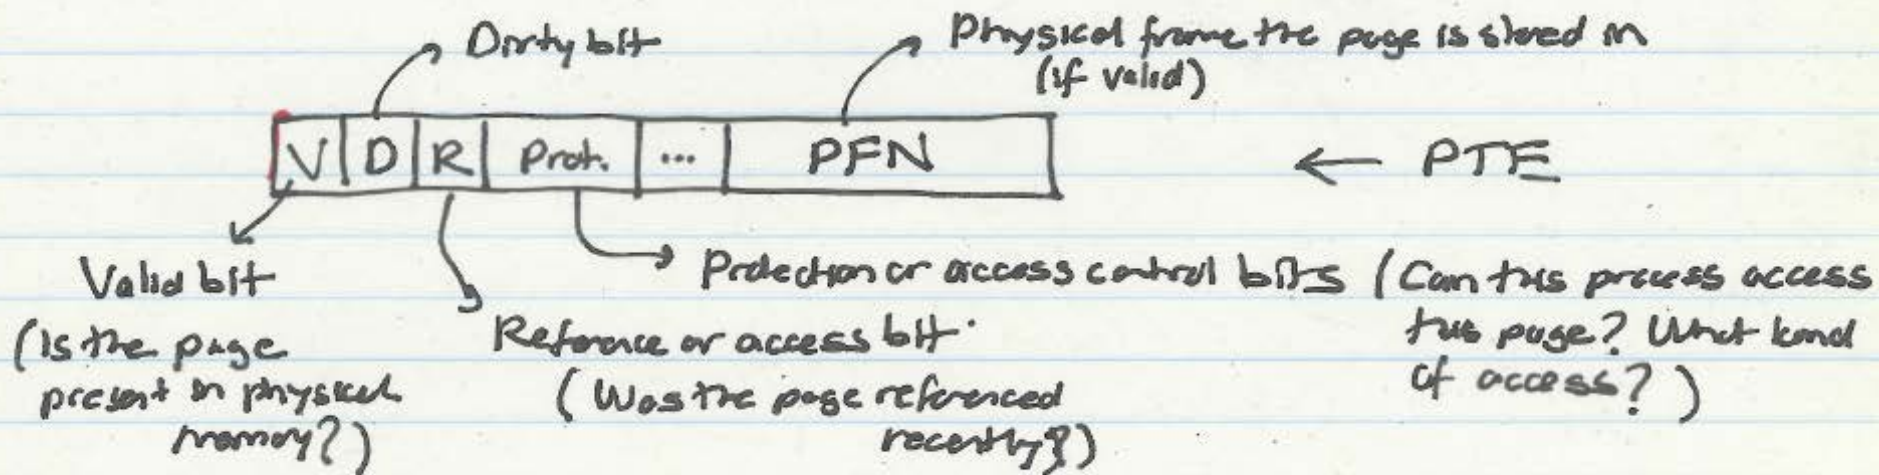
Address Translation

- Separate (set of) page table(s) per process
- VPN forms index into page table (points to a page table entry)
- Page Table Entry (PTE) provides information about page



What Is in a Page Table Entry (PTE)?

- Need a **valid** bit → to indicate validity/presence in physical memory
- Need PPN → to support translation
- Need bits to support **replacement**
- Need a **dirty** bit to support “write back caching”
- Need **protection bits** to enable access control and protection



Some Issues in Virtual Memory

Three Major Issues


- How large is the page table and how do we store and access it?
- How can we speed up translation & access control check?
- When do we do the translation in relation to cache access?
- There are many other issues we will not cover in detail
 - What happens on a context switch?
 - How can you handle multiple page sizes?
 - ...

Virtual Memory Issue I

- How large is the page table?
- Where do we store it?
 - In hardware?
 - In physical memory? (Where is the PTBR?)
 - In virtual memory? (Where is the PTBR?)
- How can we store it efficiently without requiring physical memory that can store all page tables?
 - Idea: multi-level page tables
 - Only the first-level page table has to be in physical memory
 - Remaining levels are in virtual memory (but get cached in physical memory when accessed)

Issue: Page Table Size

- Suppose 64-bit VA, 40-bit PA and 4kB page size
how large is the page table?
 - 2^{52} entries * ~ 4 bytes per entry $\approx 16 * 10^{15}$ bytes
 - And that is for just one process
 - And the process may not be using the entire VM space



Solution: Multi-Level Page Tables

- Four-level paging in x86

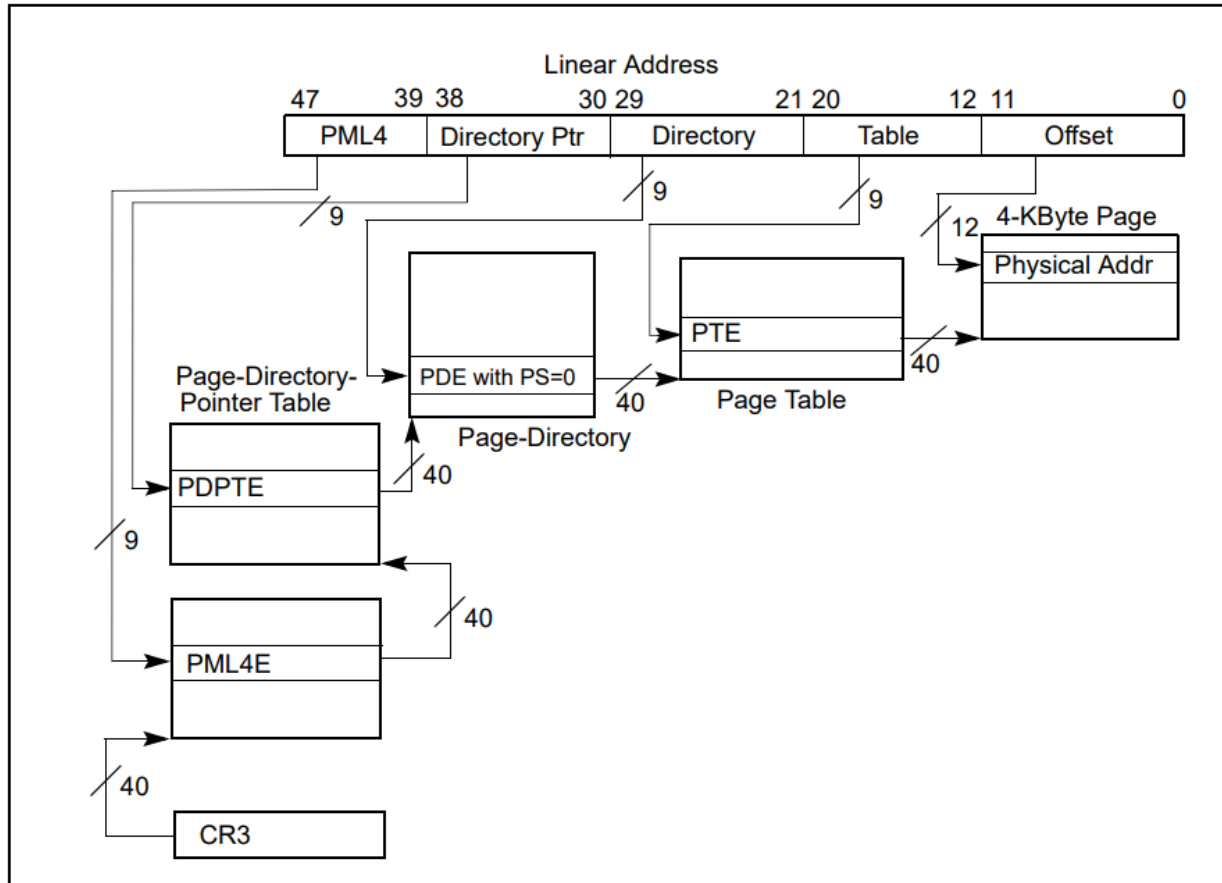


Figure 4-8. Linear-Address Translation to a 4-KByte Page using 4-Level Paging

Solution: Multi-Level Page Tables

- Four-level paging in x86

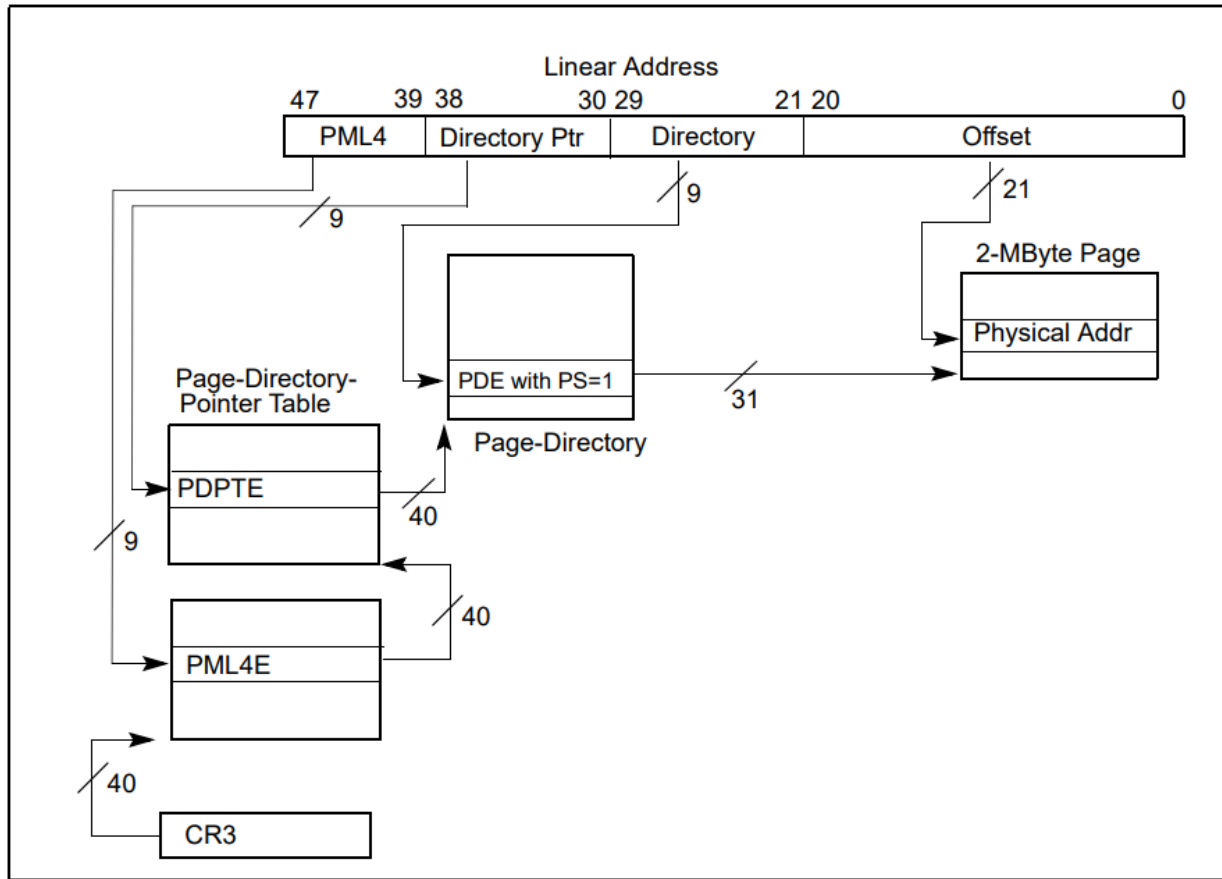



Figure 4-9. Linear-Address Translation to a 2-MByte Page using 4-Level Paging

Virtual Memory Issue II

- How fast is the address translation?
 - How can we make it fast?
- Idea: Use a hardware structure that caches PTEs → Translation lookaside buffer
- What should be done on a TLB miss?
 - What TLB entry to replace?
 - Who handles the TLB miss? HW vs. SW?
- What should be done on a page fault?
 - What virtual page to replace from physical memory?
 - Who handles the page fault? HW vs. SW?

Speeding up Translation with a TLB

- Essentially a cache of recent address translations
 - Avoids going to the page table on every reference
- Index = lower bits of VPN
- Tag = unused bits of VPN + process ID
- Data = a PTE
- Status = valid, dirty



Virtual Memory and Cache Interaction


Address Translation and Caching

- When do we do the address translation?
 - Before or after accessing the L1 cache?
- In other words, is the cache virtually addressed or physically addressed?
 - Virtual versus physical cache
- What are the issues with a virtually addressed cache?
- **Synonym problem:**
 - Two different virtual addresses can map to the same physical address
→ same physical address can be present in multiple locations in the cache → can lead to inconsistency in data


Homonyms and Synonyms

- **Homonym: Same VA can map to two different PAs**
 - Why?
 - VA is in different processes
- **Synonym: Different VAs can map to the same PA**
 - Why?
 - Different pages can share the same physical frame within or across processes
 - Reasons: shared libraries, shared data, copy-on-write pages within the same process, ...
- Do homonyms and synonyms create problems when we have a cache?
 - Is the cache virtually or physically addressed?


Cache-VM Interaction



PIPT

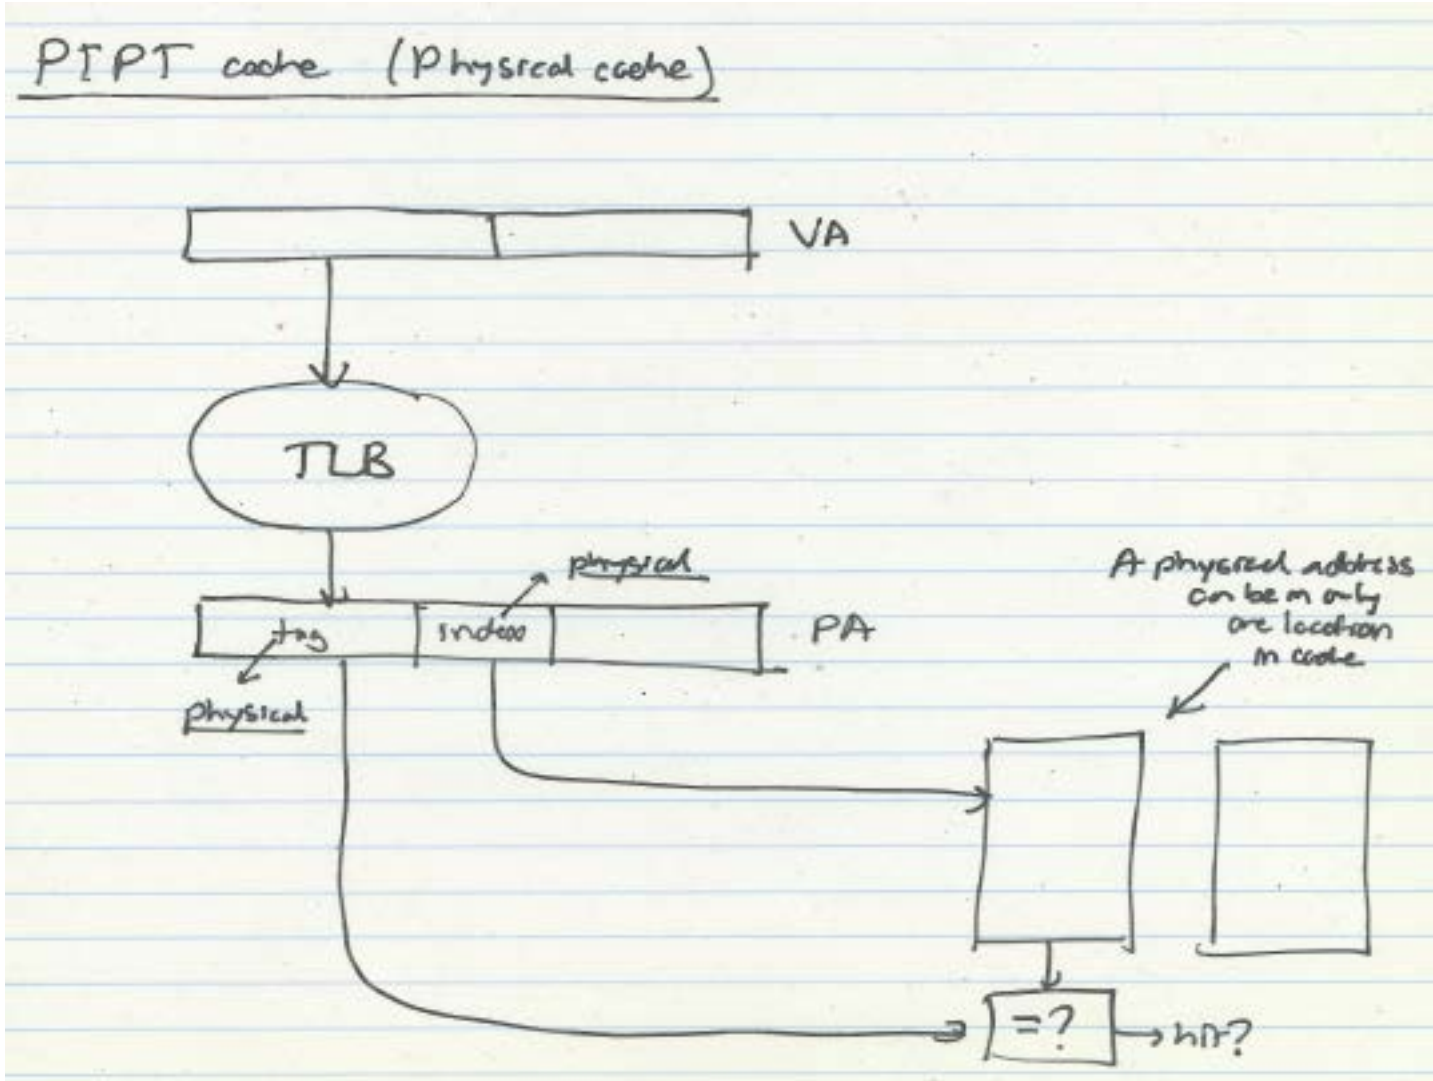


VIVT



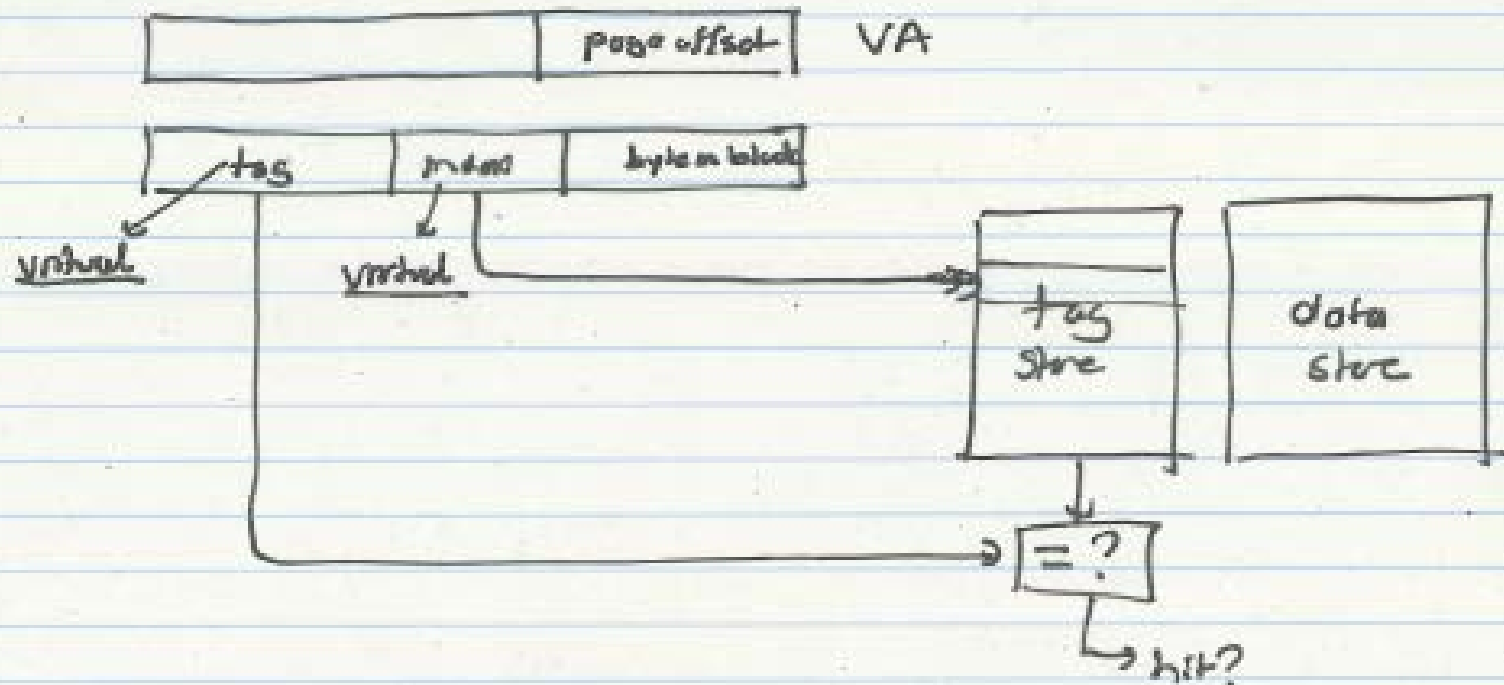
VIPT

Physical Cache




Virtual Cache

VINT cache (Virtual Cache)



VIPT

VIPT cache



Where can the same physical address be in the cache?

Slides Credit

- Attiano Purpura-Pontoniere
- Yuchen Hao
- Onur Mutlu